

Distinguishing Adult and Youth Faces Using Convolutional Neural Networks

Marta Fuentes-Filp, 2024

Introduction & Potential Applications

- Objective: Develop a CNN model to distinguish between adult and youth faces.
- Tools Used: TensorFlow, Keras

- Law Enforcement
- Social Media Platforms
- Marketing and Retail

Data Acquisition

Wernher Krutein's Photovault[®].com

Explore over 500,000 Images in my personal collection

[Click Here for Advanced searching](#)

[Animals](#) [Cities](#) [Disasters](#) [Entertainment](#) [Food](#) [Government](#) [Health](#) [Industry](#)
[Insects](#) [Military](#) [Nature](#) [People](#) [Sports](#) [Technology](#) [Universe](#) [Vehicles](#)



Your comprehensive source for

Exquisite images

royalty free and rights managed

- ✓ Many images as low as \$3.
- ✓ Over 500,000 images and growing daily.
- ✓ Advanced Search using our exclusive "FindFave" technology.

The Data: 14,635 images



PLP: 5,423 youth images



POR: 9,212 adult images

Potentially Problematic Data



Data Preparation

Keras `image_dataset_from_directory()`

- Creating training and validation sets
- Addressing class imbalance
- Data Split:
 - 80% Training Set: 11,707 images
 - 20% Validation Set: 2,926 images

CNNs & Computer Vision

- Convolutional Layer
- Pooling Layers
- Regularization
- Loss Function Optimizer

Fine Tuning the Network

Architecture

Regularization



Land of the Overfit

Land of the Underfit

CNN 4 Architecture

CNN 4: Input Layer + Multiple Hidden Layers with Dropout

Layer Type	Filters	Kernel Size	Activation	Input Shape	Additional Parameters
Conv2D	512	3	relu	(256, 256, 3)	
MaxPooling2D		2			padding='same'
Dropout					rate=0.5
Conv2D	256	3	relu		kernel_regularizer=regularizers.l2(0.03)
MaxPooling2D		2			padding='same'
Dropout					rate=0.5
Conv2D	256	3	relu		kernel_regularizer=regularizers.l2(0.03)
MaxPooling2D		2			padding='same'
Dropout					rate=0.5
Conv2D	256	3	relu		kernel_regularizer=regularizers.l2(0.03)
MaxPooling2D		2			padding='same'
Dropout					rate=0.5
Flatten					
Dense			sigmoid		Output: 1 neuron

Model Progression

CNN 1

- Good initial performance with a simple architecture.
- Lacks complexity to capture intricate patterns.

CNN 2

- Improved accuracy due to additional hidden layers
- Overfitting likely due to lack of regularization.

CNN 3

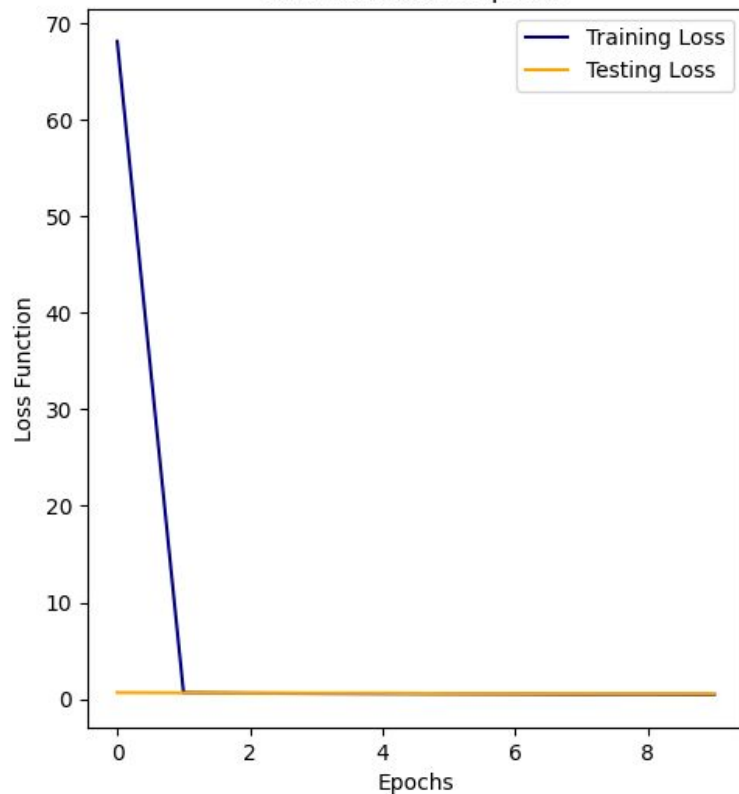
- Regularization (L2, dropout) added; performance decreased,
- Regularization may need tuning.

CNN 4


- High dropout reduced overfitting
- led to underfitting, causing lower accuracy

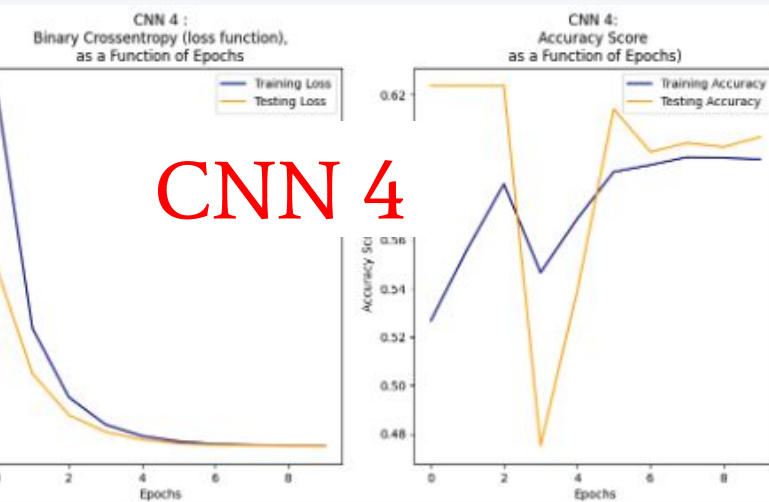
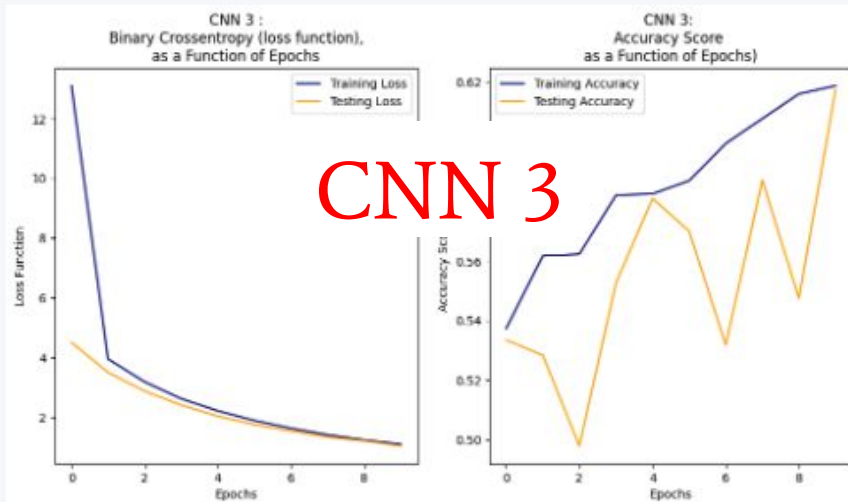
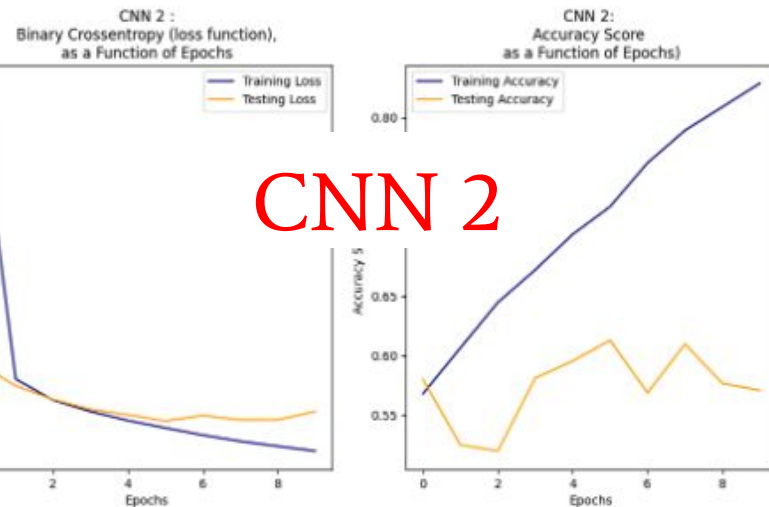

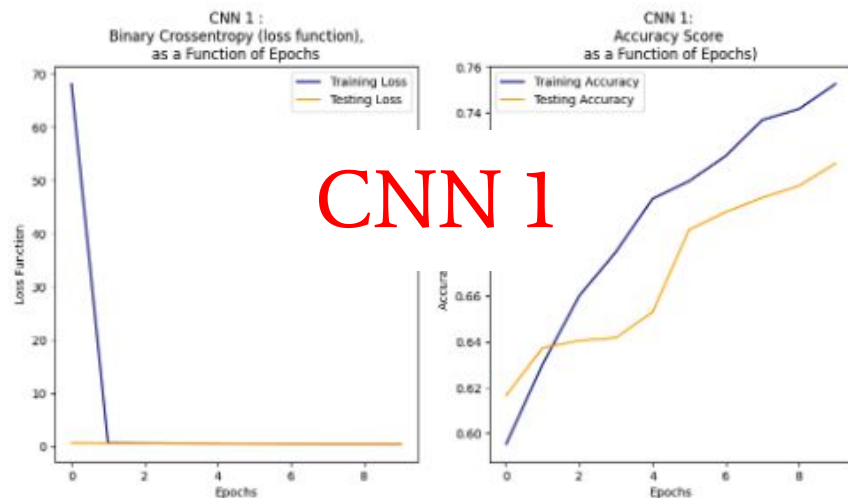
CNN1: Loss Function & Accuracy Graphs

CNN 1 :
Binary Crossentropy (loss function),
as a Function of Epochs



CNN 1:
Accuracy Score
as a Function of Epochs





Challenges

- Computational Resources
- Understanding Layers
- Inconsistent Images
- Augmentation vs. weights

Future Work

- Visualizing the images that are being misclassified
- Experimenting with ResNet
- Improving the dataset
- Augmenting data
- Explore regularization techniques
- Weed through the images
- Bounding boxes
- Find more computational power

Conclusions

- This project successfully developed and evaluated four CNN models for distinguishing between the faces of adults and youth.
- While the models demonstrated promising initial results, further improvements and optimizations are essential for real-world applications.
- Understanding the reasons behind the models' poor performance provides several key benefits:

Note: an extensive list of sources is available in the README.md file

What Questions Do You Have?